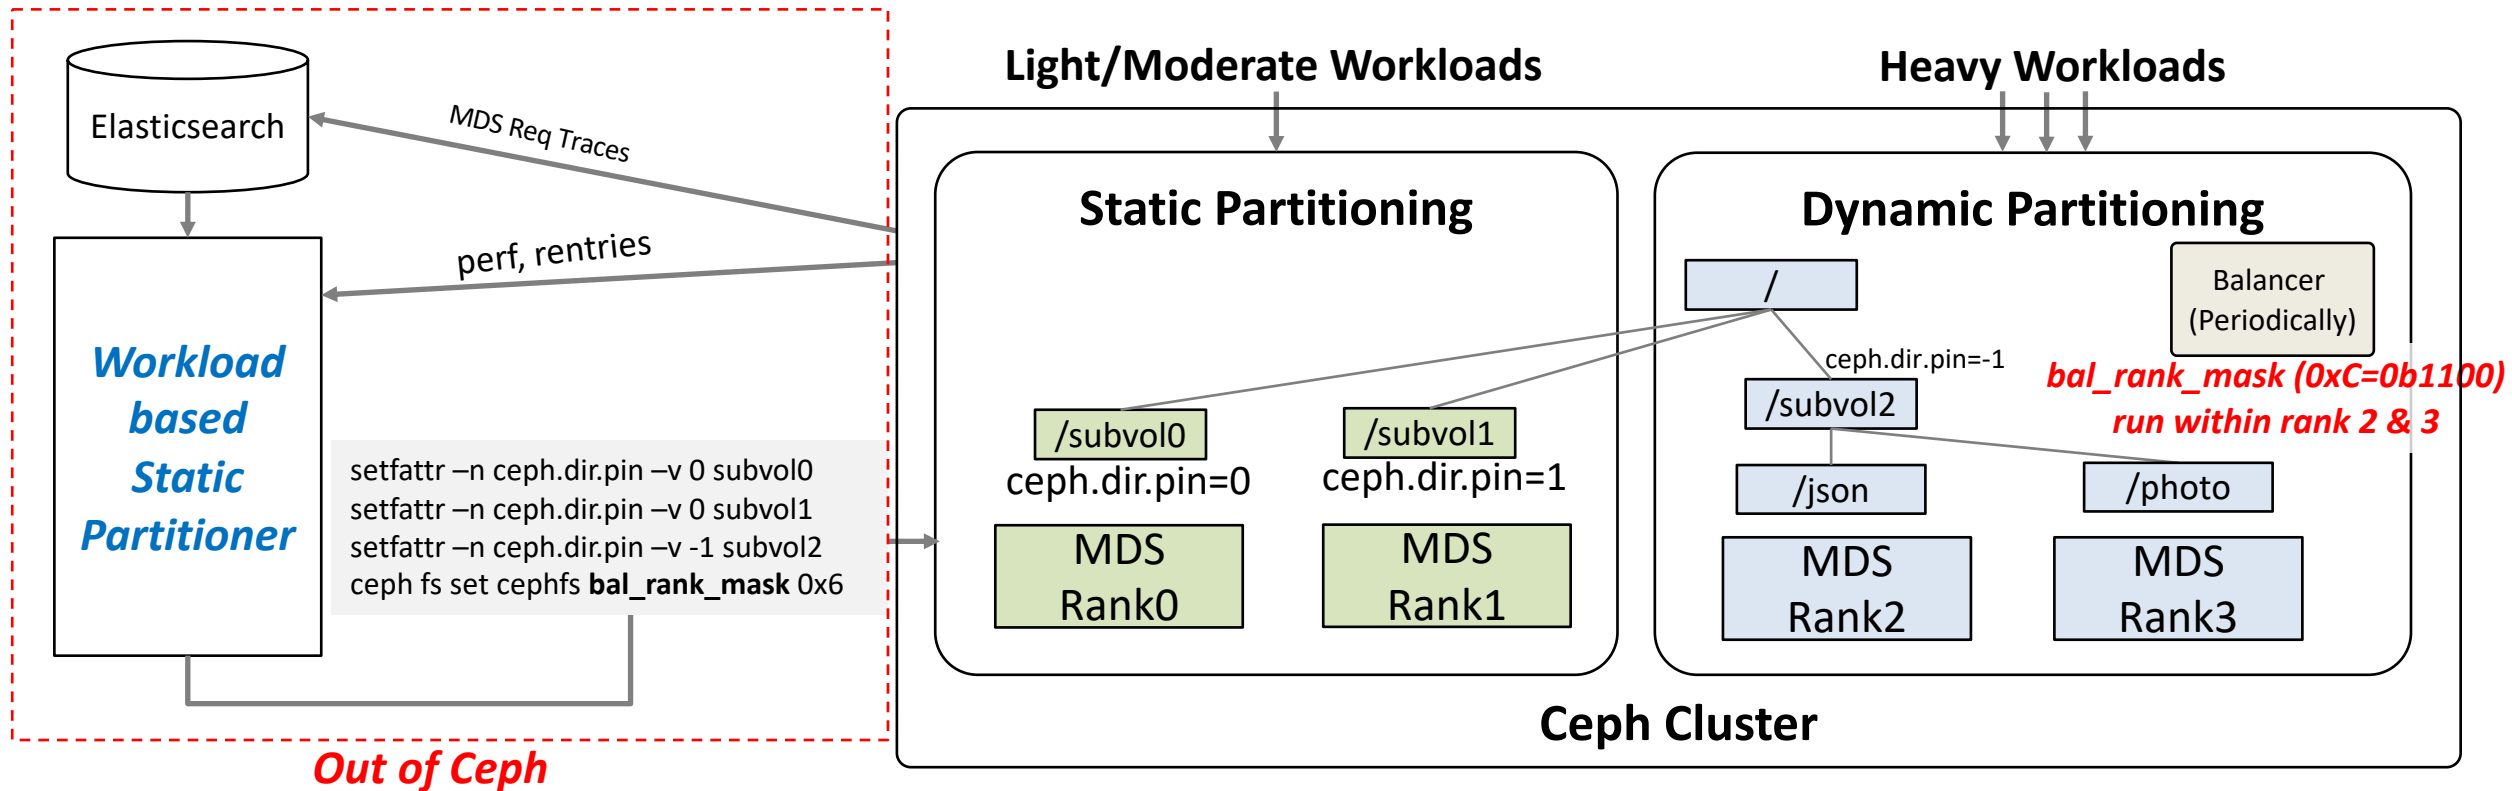


A New MDS Partitioner for CephFS

LINE
Yongseok Oh



Workload Based Static Partitioner with bal_rank_mask



- This idea was presented at **Cephalocon2023**
- **Workload based static partitioner** pins subvolumes
 - Workload calculation based on working set, reentries, and performance
 - Rarely or manually conducted if loads are uneven or latencies get higher
 - Make subvolumes involving heavy workloads managed by MDS balancer with **bal_rank_mask**

Technical Issues

- Our in-house partitioner is useful for performance
 - Compared to simple pinning, it distributes subdirs based on workloads
 - However, it is unavailable as open source
 - It needs to be revised and reimplemented for Ceph community
- `bal_rank_mask` needs to be enhanced
 - It can isolate unpinned large subtrees within certain MDS ranks from pinned subtrees
 - But, migrating large subdirs incur metadata movements
 - per subdir rank mask will be useful

A New MDS Partitioner

- rank mask option per subdir as a virtual extended attribute
 - A target subdir is dynamically within certain MDS ranks (e.g., rank0 and 1)

```
setfattr -n ceph.dir.bal.mask -v 0x3 /cephfs/home/yongseok
```

- MDS Subtree Partition Module in MGR

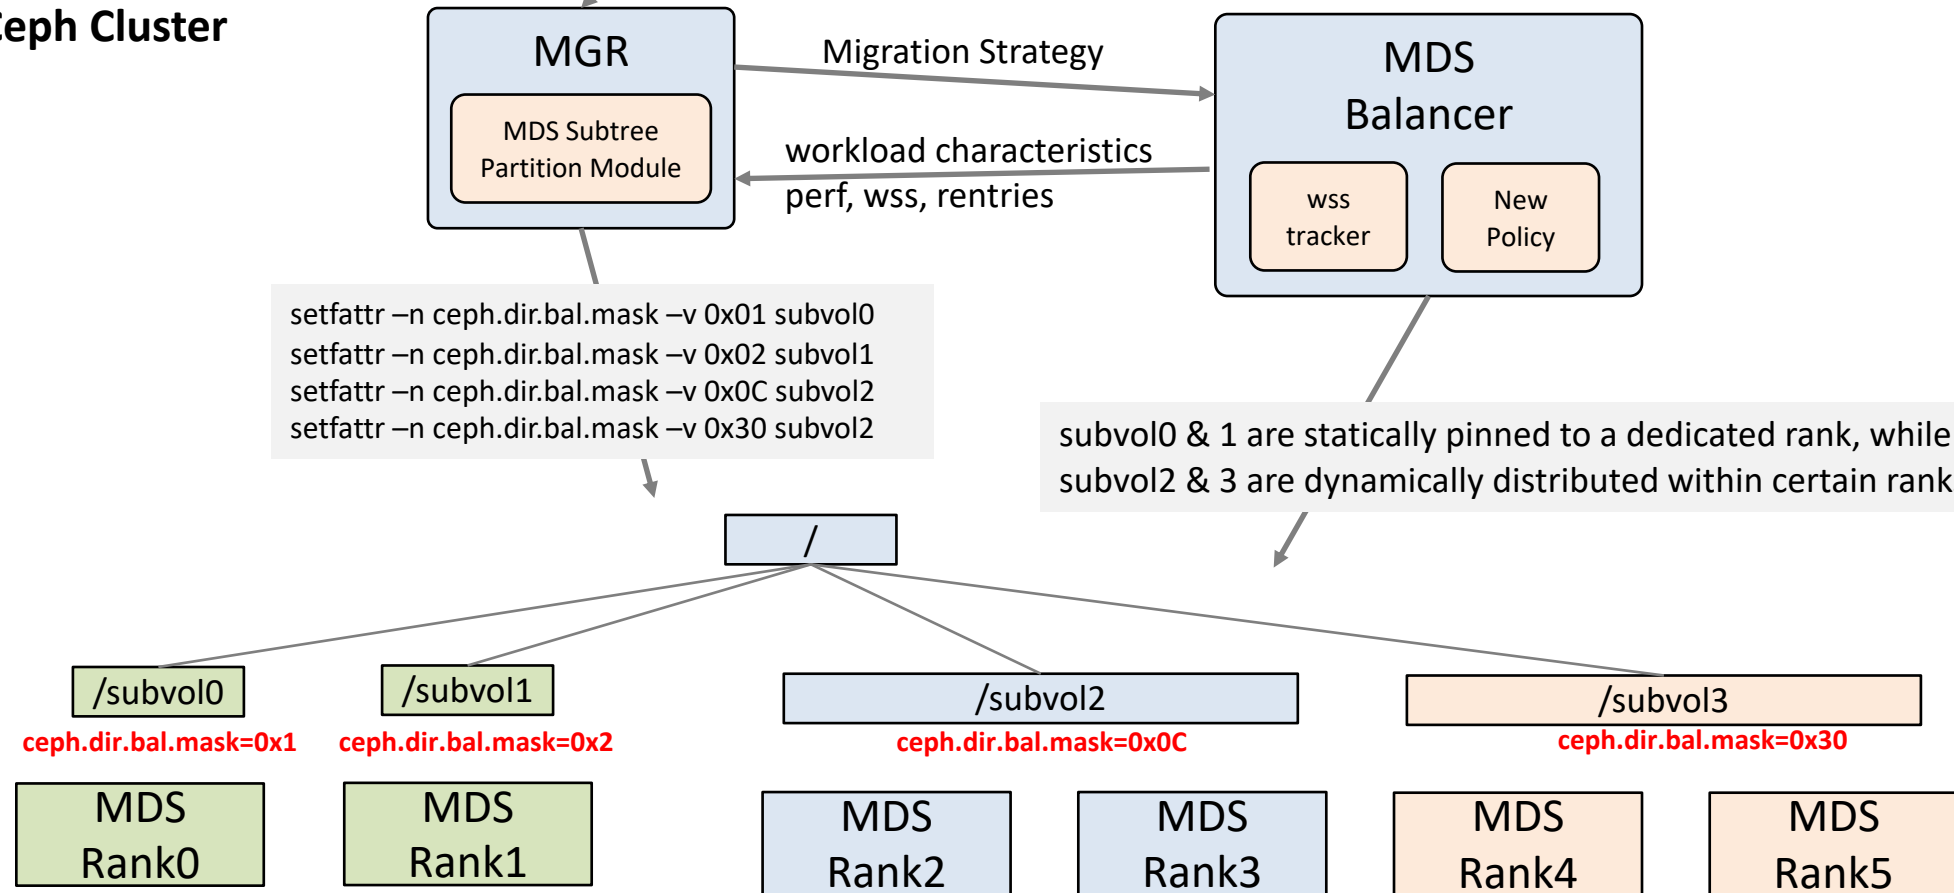
```
ceph mgr module enable mds_partitioner
ceph mds_partitioner analyze start # analyze client workloads obtained from MDSs
ceph mds_partitioner analyze status # report analysis results and recommend optimal the number of MDSs
ceph mds_partitioner partition start # start partitioning
ceph mds_partitioner partition status # report partitioning status
```

- MDS Balancer modifications
 - Working set size tracker
 - Migrate subdirs based on ceph.dir.bal.mask values of subdirs
 - Minimize MDS slow requests

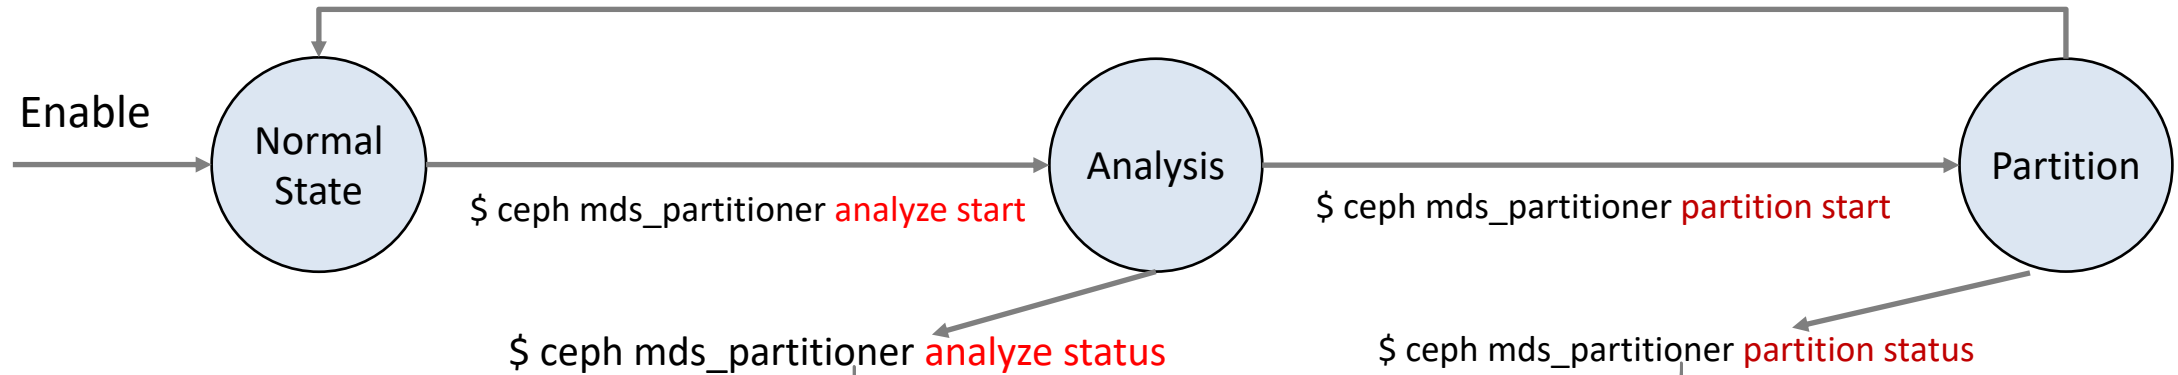
Overall Architecture

```
$ ceph mgr module enable mds_partitioner
$ ceph mds_partitioner analyze start
$ ceph mds_partitioner analyze status
$ ceph mds_partitioner partition start
$ ceph mds_partitioner partition status
```

Ceph Cluster



Example of Operation Flow



Name	wss	reqs	retries	workload	Current Ranks	New Ranks	Migration Progress	Migration Status
Subvol1	1,000,000	1,203,030	50,000,000	339	0	0,1	20%	In-progress
Subvol2	700,000	500,000	1,000,000	137	1	2	0%	Ready
subvol3	100,000	5,000	5,000,000	21	2	2	100%	Done
subvol4	3,000	20,000	70,000	3	2	2	100%	Done
Total	1,803,000	1,728,030	56,070,000	500				

wss: working set size
reqs: requests
retries: files + dirs

workload = (working_set_size / total_working_set * 2
+ requests / total_requests * 2
+ retries / total_retries) * 100