

# A New Partitioning Scheme for CephFS

LINE  
오용석



Ceph Days Korea 2023

# Outline of Contents

- Introduction and background
- CephFS subtree partitioning
- Combining static and dynamic partitioning schemes
  - Workload based static partitioner with `bal_rank_mask`
  - A new CephFS MDS balancer with `ceph.dir.bal.mask`
- Future works and conclusions


# Who Am I?



- Yongseok Oh (오용석)
- 2015 Ph.D. in Computer Science
  - Flash based storage systems
  - Research papers in FAST, MSST, Systor (over 400 citations)
- 2015 SKT - SW-defined Storage Lab.
  - Enterprise NVMe SSD FW
  - SPDK based file system
- 2017 SK hynix - FW Group
  - Enterprise NVMe SSD FW
  - ZNS, global wear leveling, SR-IOV, NVMe-MI
- 2020 ~ LINE Plus - Cloud Storage
  - CephFS based shared file service
  - NVMeoF based storage service
- Hobby: camping, boxing, jogging

# Cephalocon2023


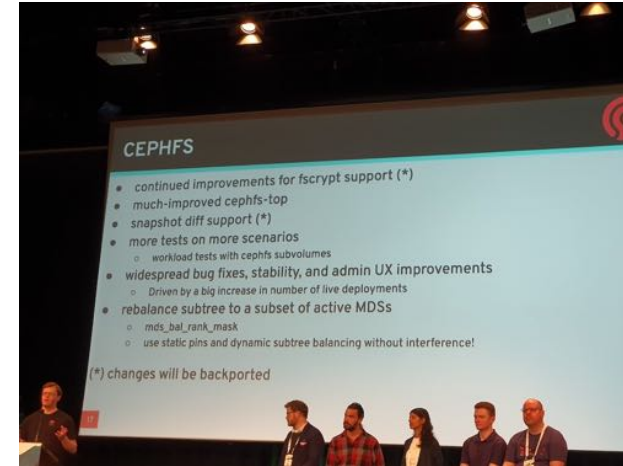
- Optimizing CephFS with Combining MDS QoS Scheduling and Static-Dynamic Subtree Partitioning



<https://www.youtube.com/watch?v=pDURll6Y-Ug>

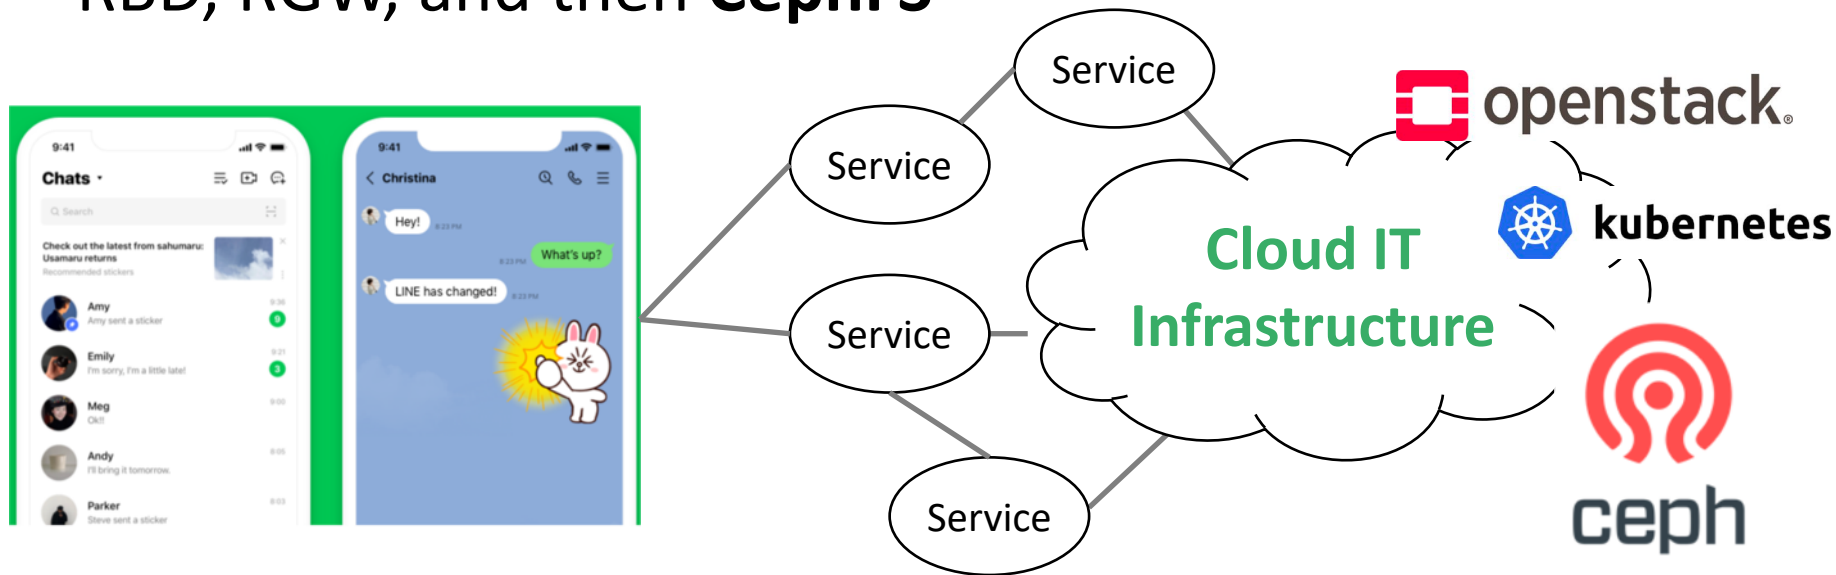


# Cephalocon2023



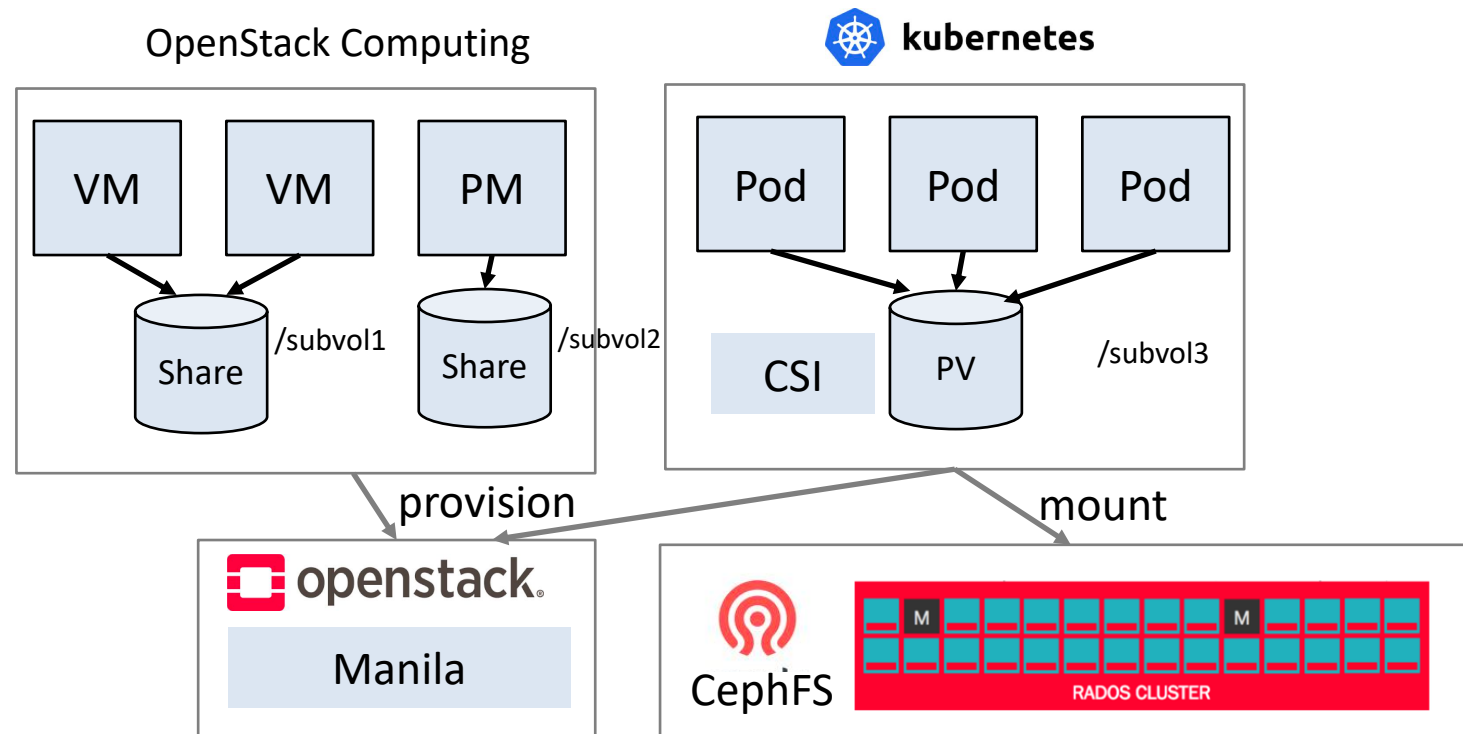
# LINE

- LINE, a messaging and communication service with over **150 million** daily active users
- OpenStack and Kubernetes as a cloud Infrastructure
- Ceph as a software defined storage
  - RBD, RGW, and then **CephFS**



# Use Cases of CephFS in LINE

- CephFS as a shared file system for clouds
  - OpenStack Manila to manage subvolumes of CephFS
- File system can be shared to VMs/Pods (major benefit)
  - NAS, ML training, speech data, backup





# Key Features of CephFS

- POSIX compliant distributed file system
- Multi active MDSs (e.g., scalability)
- Dynamic/static sub directory partitioning
- Standby/standby-replay MDSs (e.g., rapid failover)
- Journaling (e.g., guaranteeing metadata consistency)
- Kernel/FUSE/libcephfs client support
- Subvolume/snapshot/quota management
- QoS (not support)



# CephFS Architecture



# CephFS Clusters in LINE

- Major two clusters of our environment

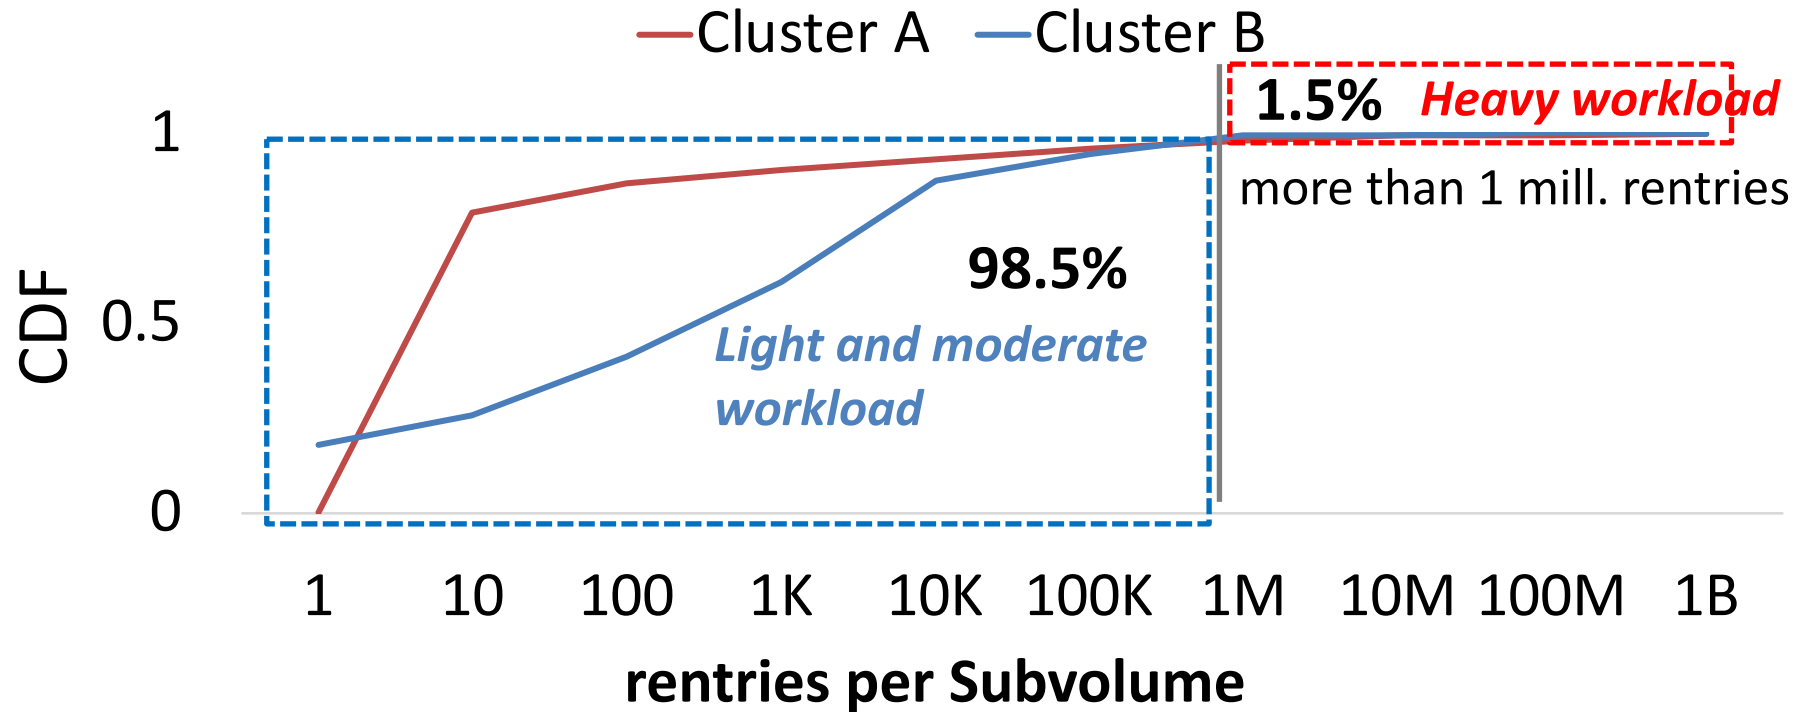
	Cluster A	Cluster B
<b>Workload</b>	General Purpose	ML, Data Processing
<b>Active MDSs (on PMs + VMs)</b>	30	36
<b>OSDs (SSD based)</b>	PB scale	PB scale
<b>Sessions</b>	> 5,000	> 5,000
<b>Subvolumes</b>	> 3,000	> 3,000

# On Deploying CephFS in Production

- Sep 2020: CephFS service with two active MDSs + standby
- Jan 2021: technical issues and outages with growing users
  - Noisy neighbor, MDS slow request by balancer, increased recovery time
- Feb 2021: simple static partitioning has been applied
- Sep 2021: proposed MDS QoS scheduler (e.g., [PR 38506](#))
- Dec 2021: enhanced MDS recovery time (e.g., [PR 44246](#) , [PR 41358](#))
  - Becoming up:active took 2 hours because of heartbeat timed out
- Aug 2022: the number of MDSs has increased to up to 36

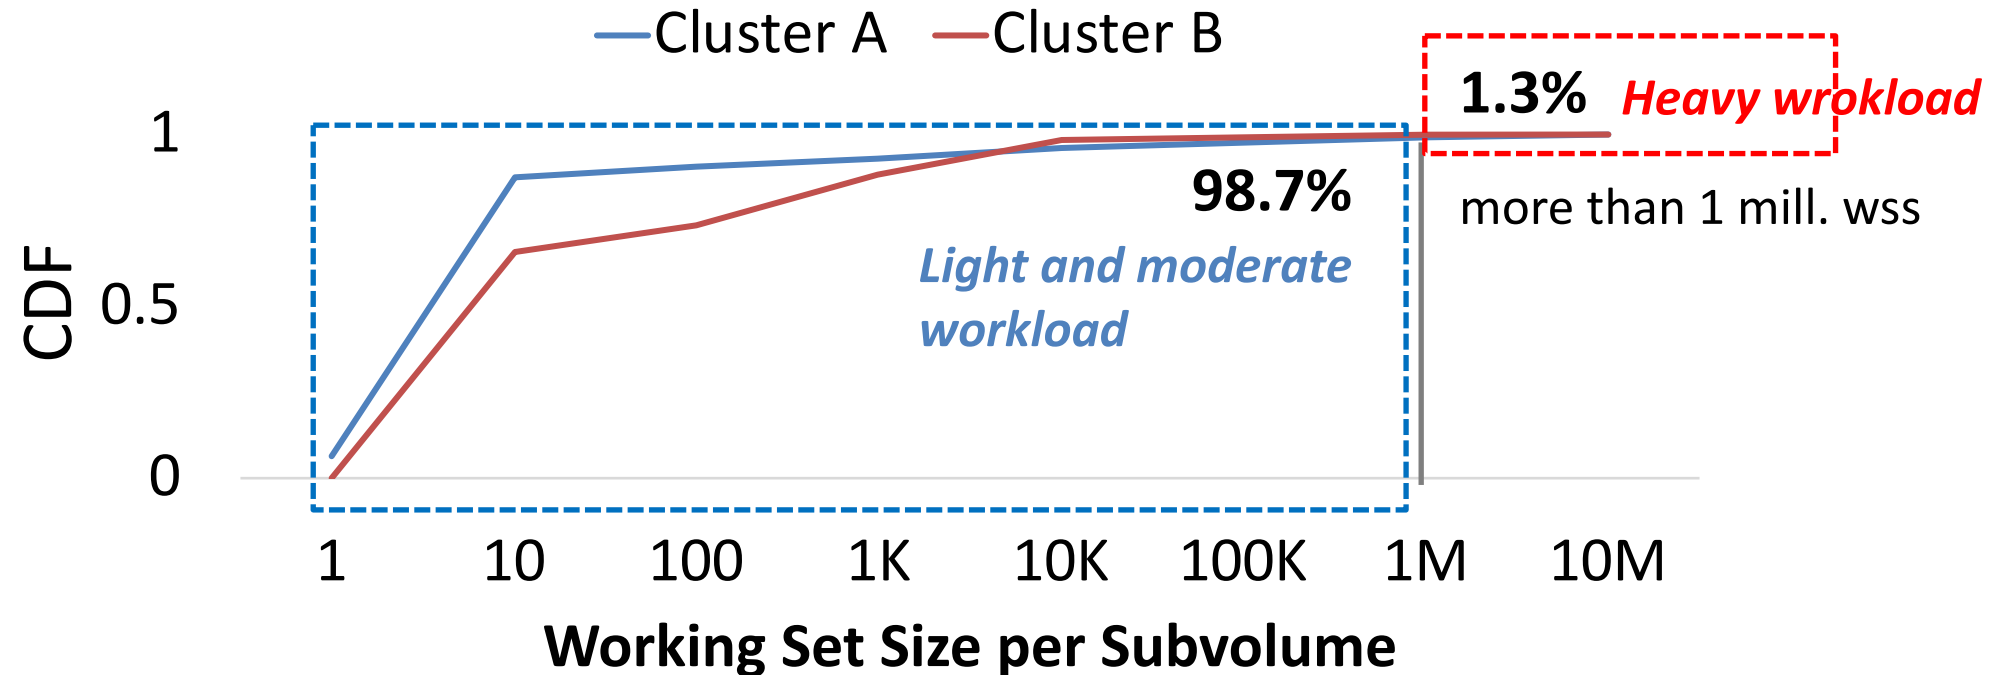
Now, we are working on ***subtree partitioning!***

# MDS Workloads: rentries Distribution



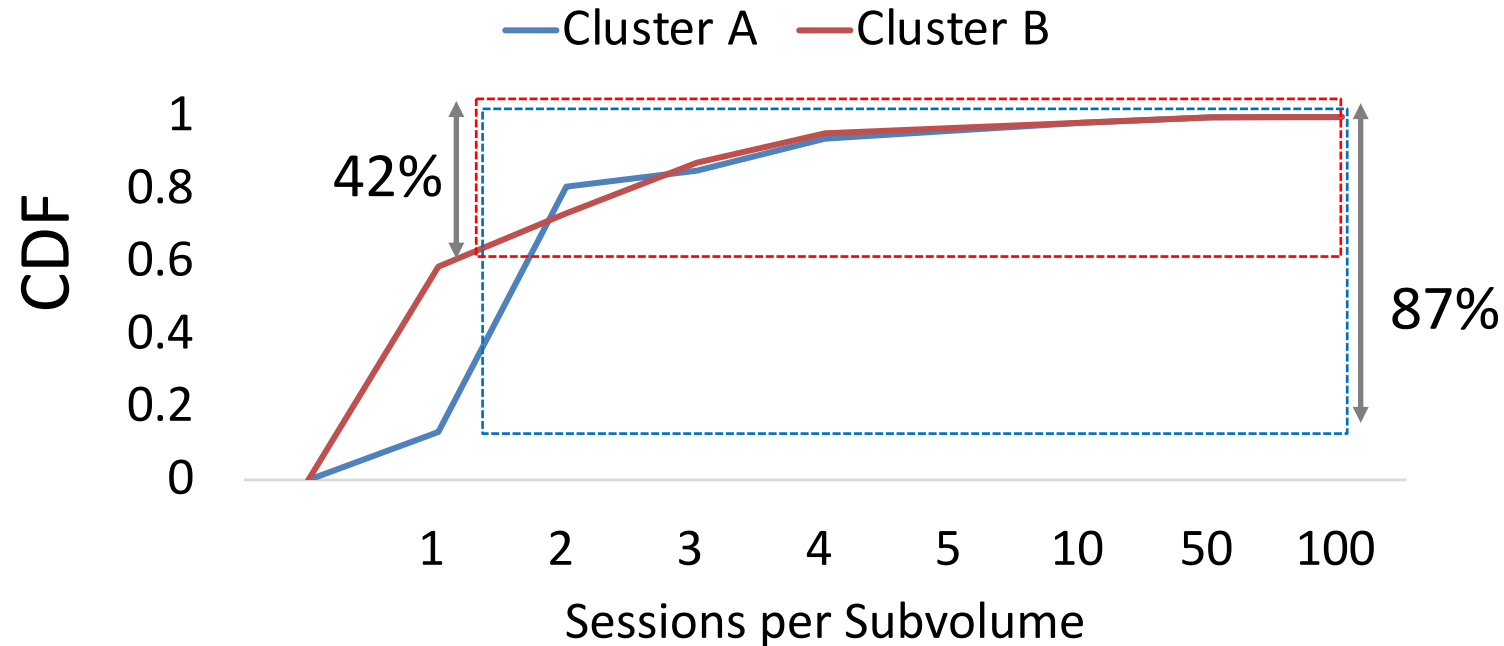
- **rentries** (e.g., files + dirs) per subvolume (measured via ceph.dir.rentries)
  - 98.5% of subvolumes contain rentries less than or equal to 1 mil.
  - Remaining of subvolumes have more than 1 mil rentries
- **Heavy workload** may affect the cache hit ratio, leading to performance degradation

# MDS Workloads: Working Set Distribution



- **Working set size (WSS)** per subvolume (collected via `uniq()` of Elasticsearch)
  - 98.7% subvolumes have WSS of less than 1 mil.
- Heavy workload requires the large cache memory

# MDS Workloads: Sessions per Subvolume



- 42% and 87% subvolumes are shared with sessions that have at least two sessions
  - Many sessions have the potential to generate heavy workloads

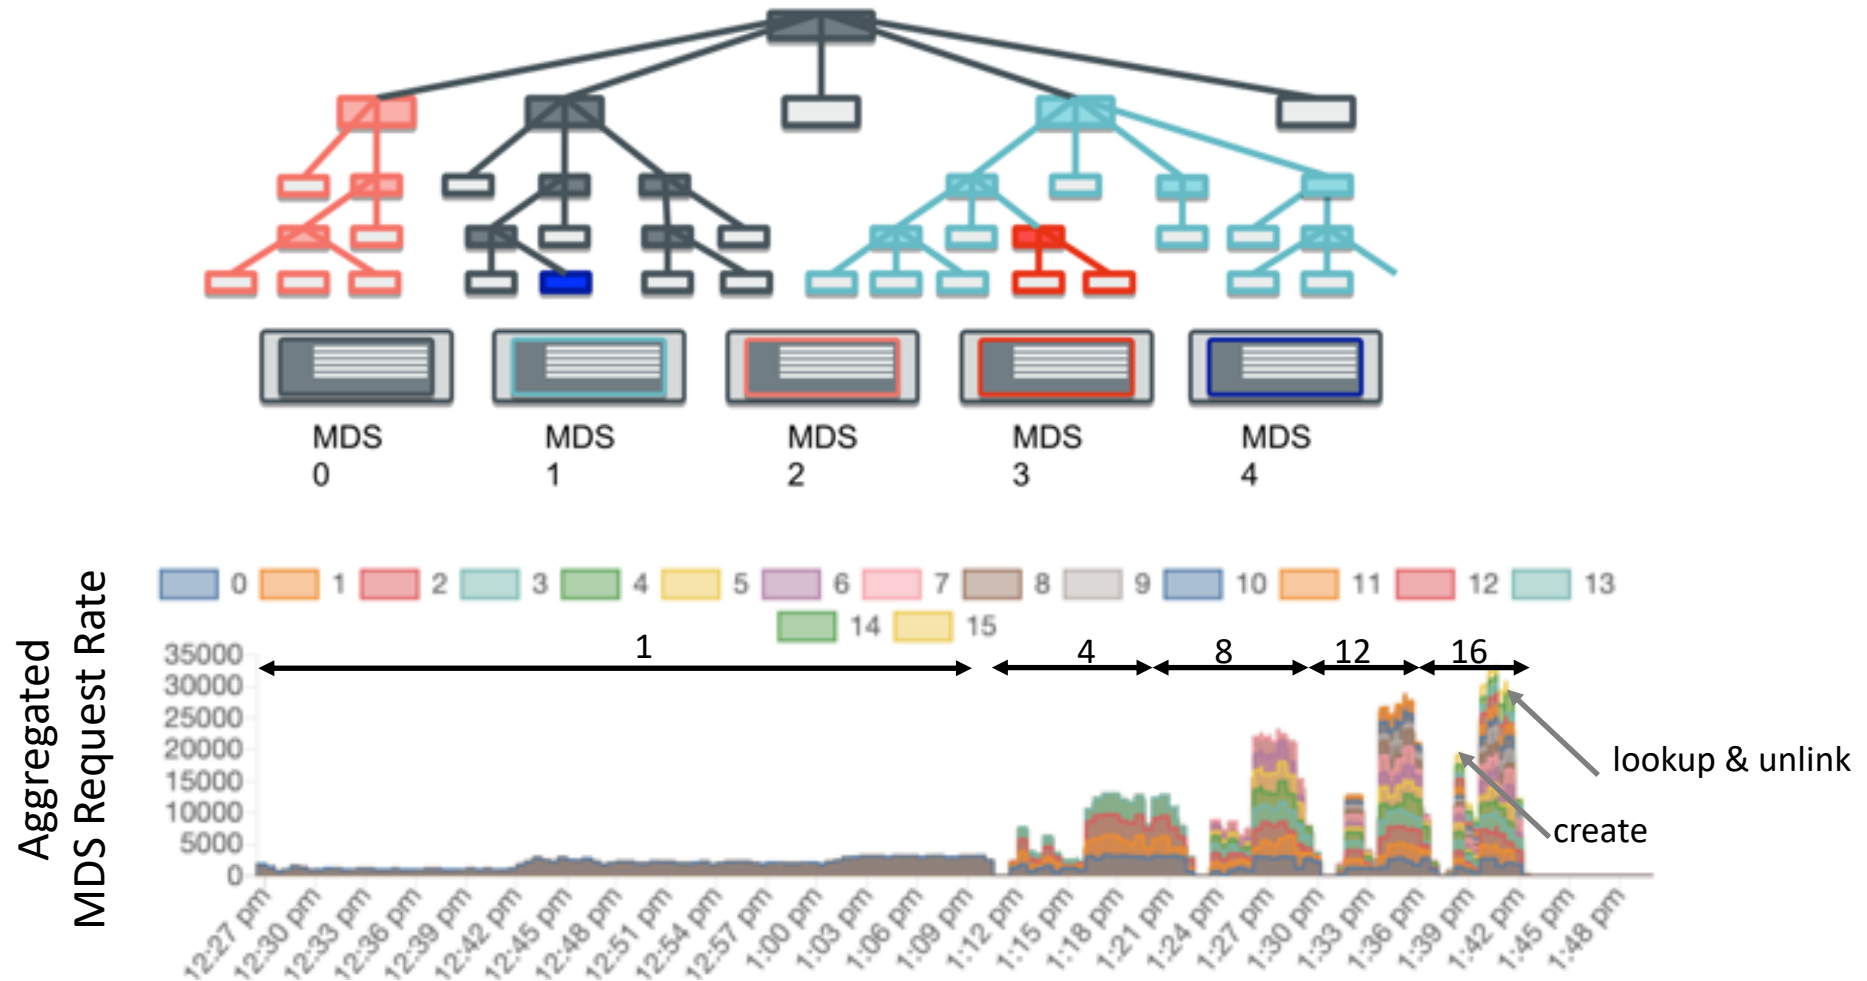


# Outline of Contents

- Introduction and background
- **CephFS subtree partitioning**
- Combining static and dynamic partitioning schemes
  - Workload based static partitioner with `bal_rank_mask`
  - A new CephFS MDS balancer with `ceph.dir.bal.mask`
- Future works and conclusions

# MDS Scalability with Subtree Partitioning

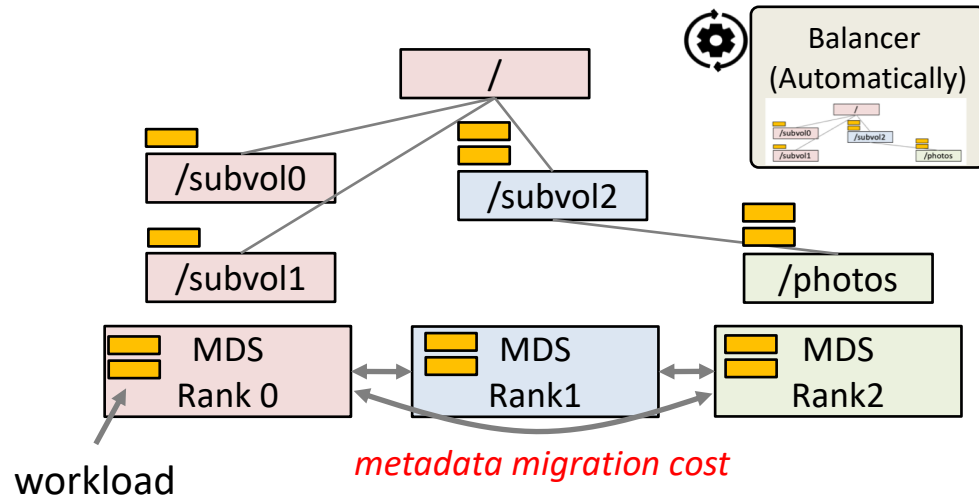
Can MDS performance scale with multiple MDSs?



# Dynamic vs Static Partitioning

## Dynamic Subtree Partitioning

Periodically redistribute subtrees



Pros

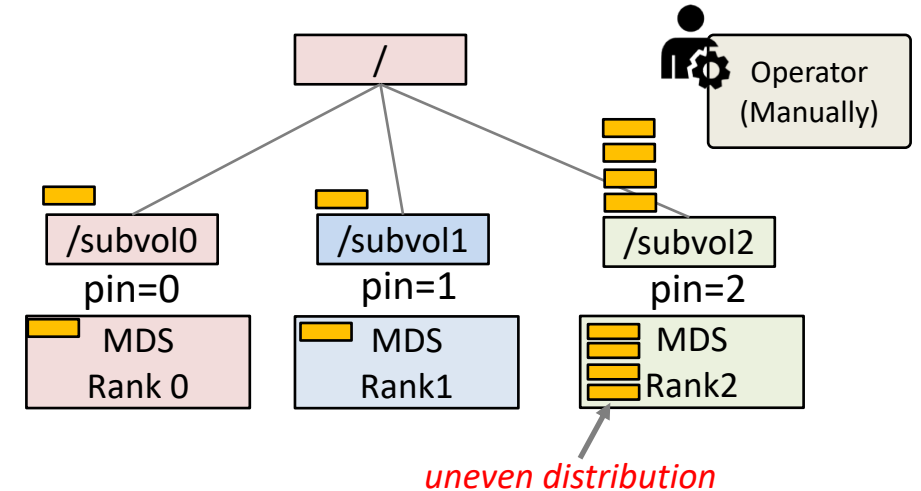
- MDS horizontal scalability

Cons

- Metadata migration cost

## Static Subtree Partitioning

Simply pin subtrees to their own ranks



Pros



- Negligible migration cost
- Provisioned performance

Cons

- Uneven workload distribution
- Additional operator efforts

# Negative Impact of Dynamic Partitioning

- Increased subtree changes and inode migrations



# Migrating Subtrees Incurs MDS Slow Requests

```
"MDS_SLOW_REQUEST": {  
  "severity": "HEALTH_WARN",  
  "summary": {  
    "message": "2 MDSs report slow requests"  
  },  
  "detail": [  
    {  
      "message": "mds(mds.1): 1253 slow requests are blocked > 30 secs"  
    },  
    {  
      "message": "mds(mds.0): 1 slow requests are blocked > 30 secs"  
    }  
  ]  
}
```

We have observed slow requests are resolved since employing static partitioning!

# Outline of Contents

- Introduction and background
- CephFS subtree partitioning
- **Combining static and dynamic partitioning schemes**
  - Workload based static partitioner with `bal_rank_mask`
  - A new CephFS MDS balancer with `ceph.dir.bal.mask`
- Future works and conclusions




# Our Idea: Combine Two Approaches

Static partitioning for moderate workloads, while dynamic partitioning for heavy workloads

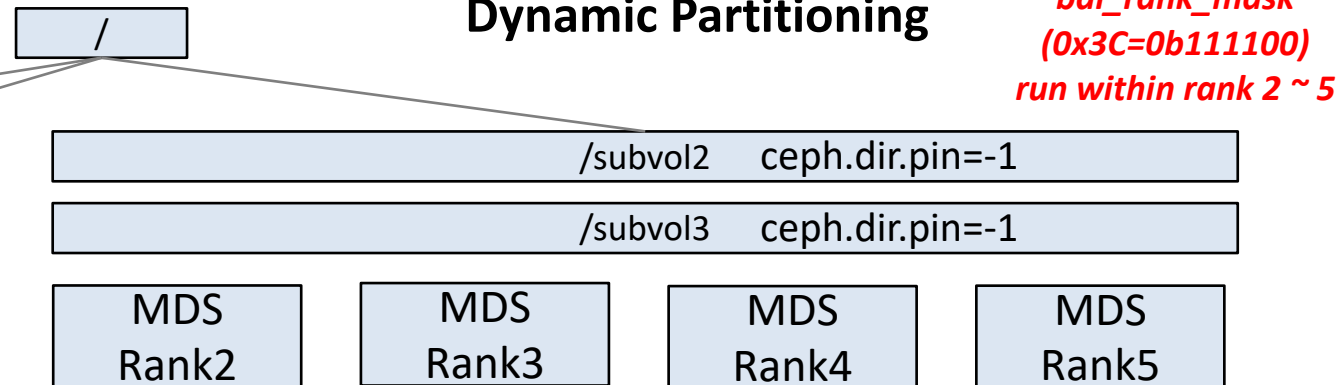
Light/Moderate Workloads

Heavy Workloads

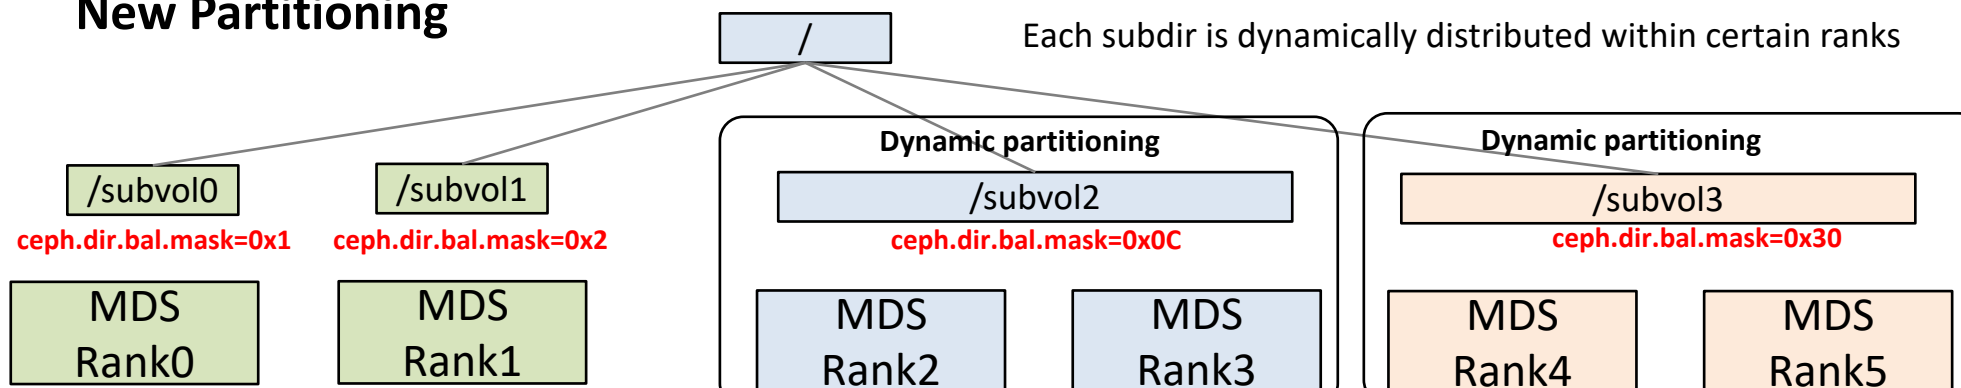
## Static Partitioning



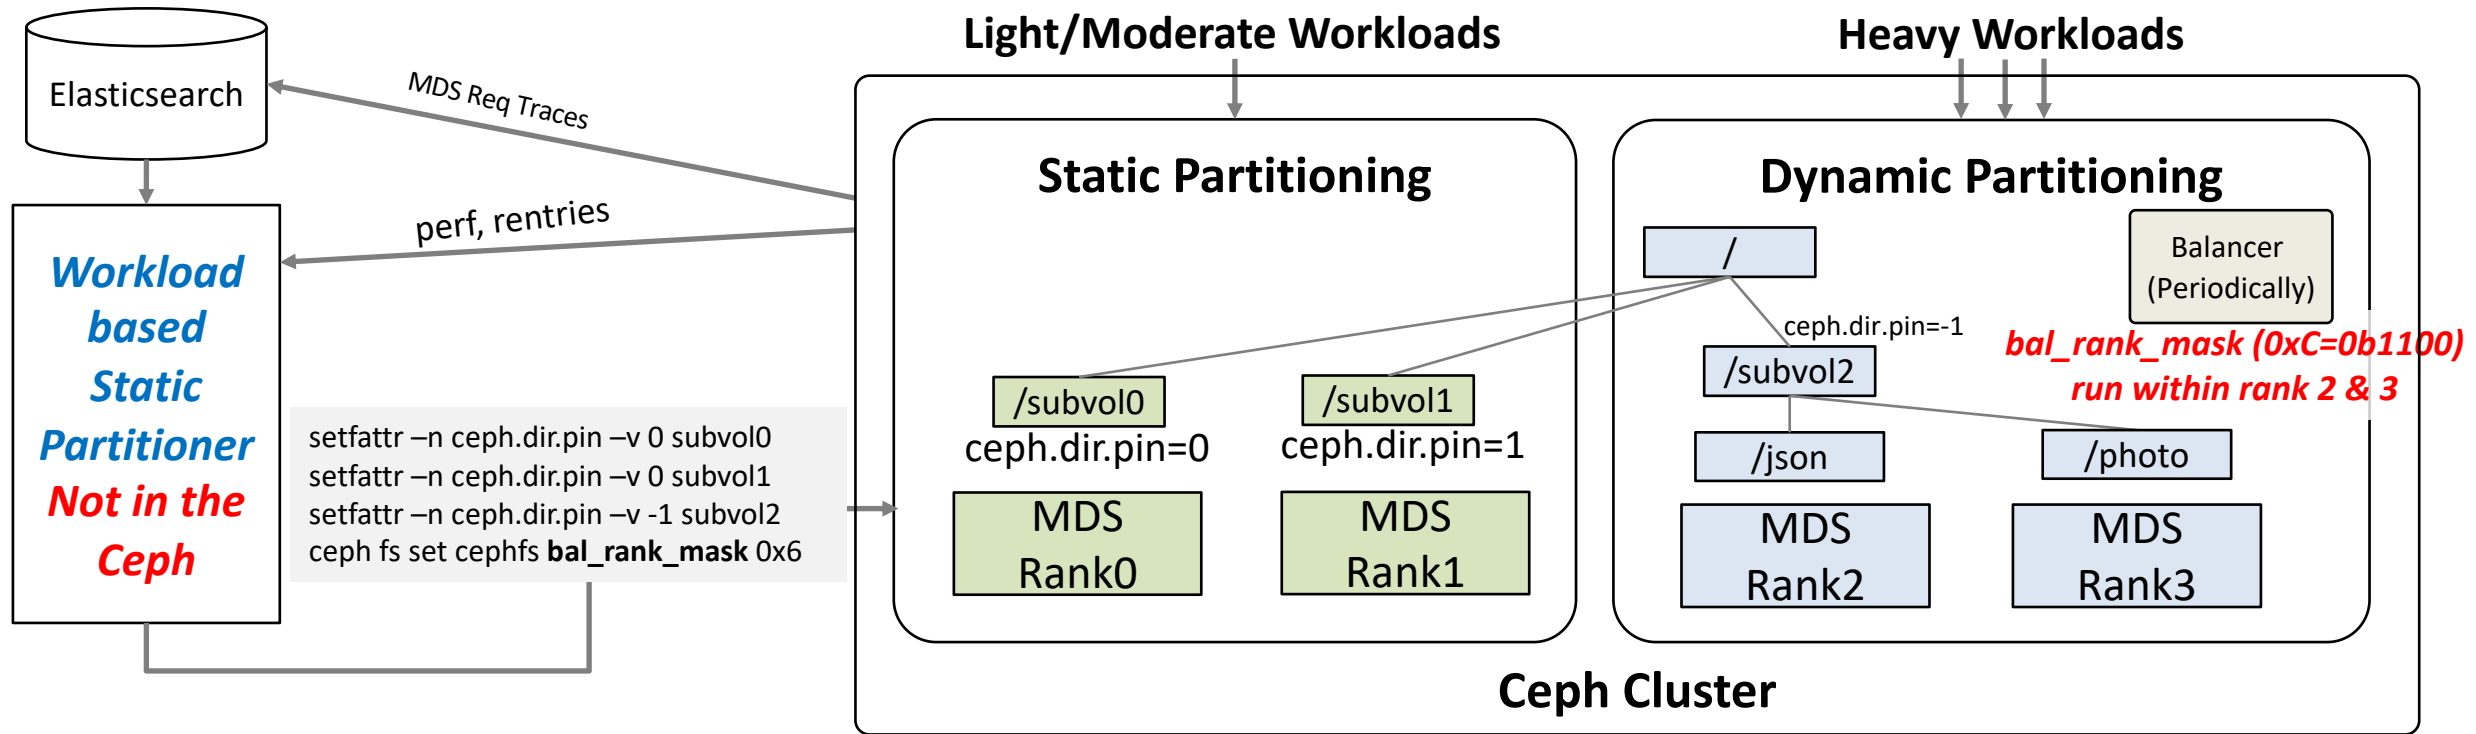
## Dynamic Partitioning



## New Partitioning



# Workload Based Static Partitioner with bal\_rank\_mask



- **Workload based static partitioner** pins subvolumes
  - Workload calculation based on working set, reentries, and performance
  - Rarely or manually conducted if loads are uneven or latencies get higher
- **bal\_rank\_mask** enables the balancer to dynamically rebalance unpinned subtrees within particular active MDS ranks [PR 43284](#)

# How to use bal\_rank\_mask

Change from 0b**111000** to 0b**000111**

```
$ ceph fs set cephfs bal_rank_mask 0x38
setting the metadata balancer rank mask to 0x38
$ sleep 180
$ ceph fs set cephfs bal_rank_mask 0x7
setting the metadata balancer rank mask to 0x7
```

```
$ ceph fs status
cephfs - 8 clients
```

=====

RANK	STATE	MDS	ACTIVITY	DNS	INOS	DIRS	CAPS
0	active	mds001	Reqs: 0 /s	5259	5095	84	14
1	active	mds017	Reqs: 0 /s	463k	463k	31.1k	21
2	active	mds018	Reqs: 0 /s	295k	295k	18.3k	38.3k
3	active	mds016	Reqs: 2 /s	53.9k	53.9k	3573	53.6k
4	active	mds015	Reqs: 405 /s	65.4k	65.5k	4018	65.0k
5	active	mds010	Reqs: 398 /s	81.9k	81.9k	4906	81.8k



```
$ ceph fs status
cephfs - 8 clients
```

=====

RANK	STATE	MDS	ACTIVITY	DNS	INOS	DIRS	CAPS
0	active	mds001	Reqs: 364 /s	75.8k	75.4k	5943	70.9k
1	active	mds017	Reqs: 259 /s	588k	587k	39.7k	123k
2	active	mds018	Reqs: 366 /s	395k	394k	26.5k	137k
3	active	mds016	Reqs: 5 /s	678	683	637	36
4	active	mds015	Reqs: 8 /s	504	509	500	83
5	active	mds010	Reqs: 4 /s	970	975	831	84

# Evaluation Environment

- Ceph Pacific 16.2.10 (integrated with our PRs) installed in Rocky8
  - [bal\\_rank\\_mask - PR 43284](#) (merged)
  - [MDS QoS Scheduler - PR 38506](#) (ready to review)
- VDBench tool generates workloads
  - Each client has its own subvolume (e.g., /volumes/\_nogroup/\$subvol)

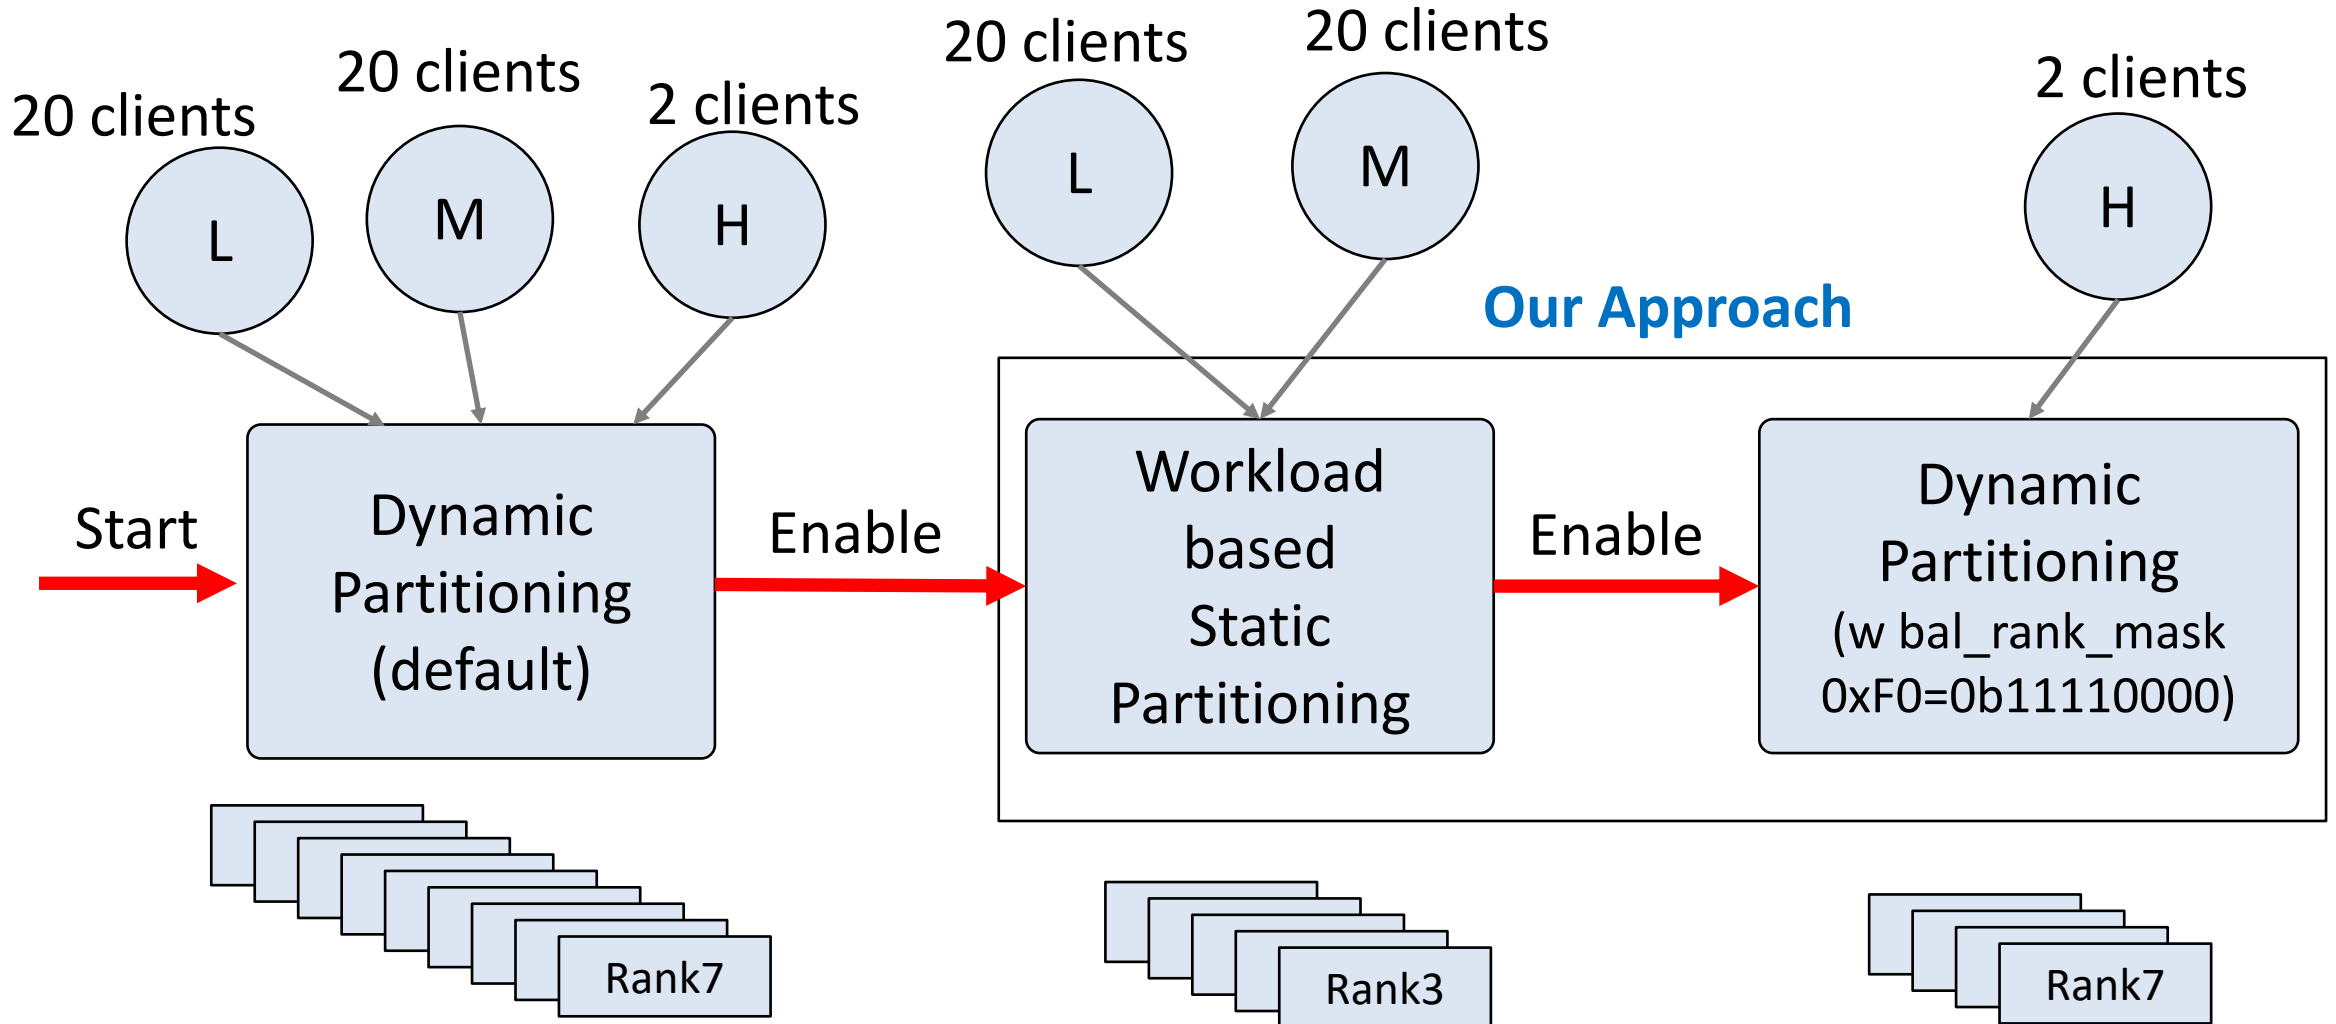
Workload Spec.

Type	Files	File Size	Clients	Threads per Client
<b>Light</b>	50K	4K	20	1
<b>Moderate</b>	500K	4K	20	1
<b>Heavy</b>	5,000K	4K	2	16

HW Spec.

Type	Server Spec	Count
MDS	4 * vCPU 32GB RAM	active 8 (mds_cache_memory_limt: 12GB)
OSD	40 * pCPU 128GB RAM	6 Servers * 6 SATA SSDs = 36
Mon	2 * vCPU 4GB RAM	3
Mgr	4 * vCPU 8GB RAM	3
Client	4 * vCPU 8GB RAM	42
Network	10Gbps	

# Evaluation Sequences




# Evaluation Results (1/3)

Balancer default

Our  
Static  
Partitioning

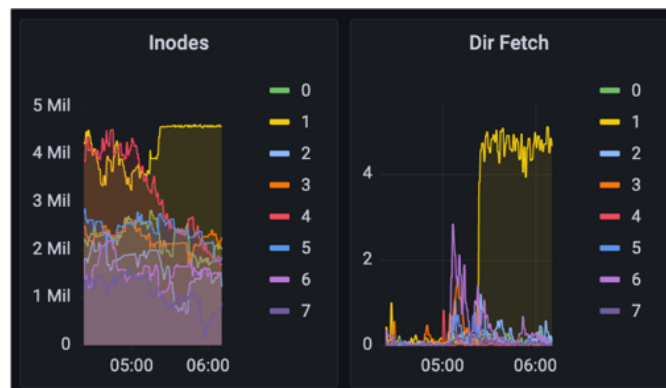
Dynamic Partitioning  
bal\_rank\_mask 0xf0




1 Performances of MDSs are fluctuated due to exporting/importing inodes



2 inodes are unevenly distributed rank1 keeps inodes more than 4mil.



Reply latencies highly increased






# Evaluation Results (2/3)

Balancer default


Our  
Static  
Partitioning

Dynamic Partitioning  
bal\_rank\_mask 0xf0



- 3 subvolumes of light/moderate workloads are moved from all ranks to rank0~3

Performance is still not recovered as rank 1 has a lot of inodes




# Evaluation Results (3/3)

Balancer default

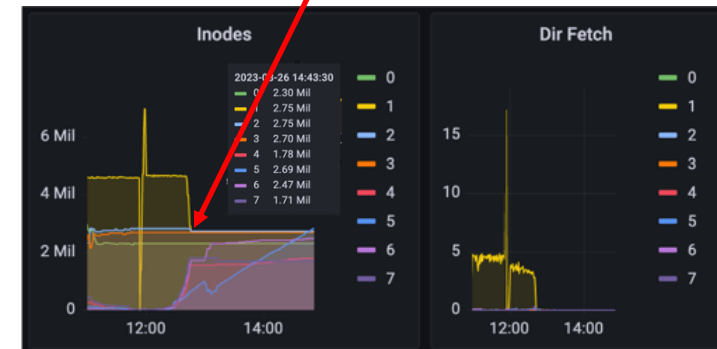
Our  
Static  
Partitioning

Dynamic Partitioning  
*bal\_rank\_mask 0xf0*



4 subvolumes of heavy workloads are migrated to rank 4~7

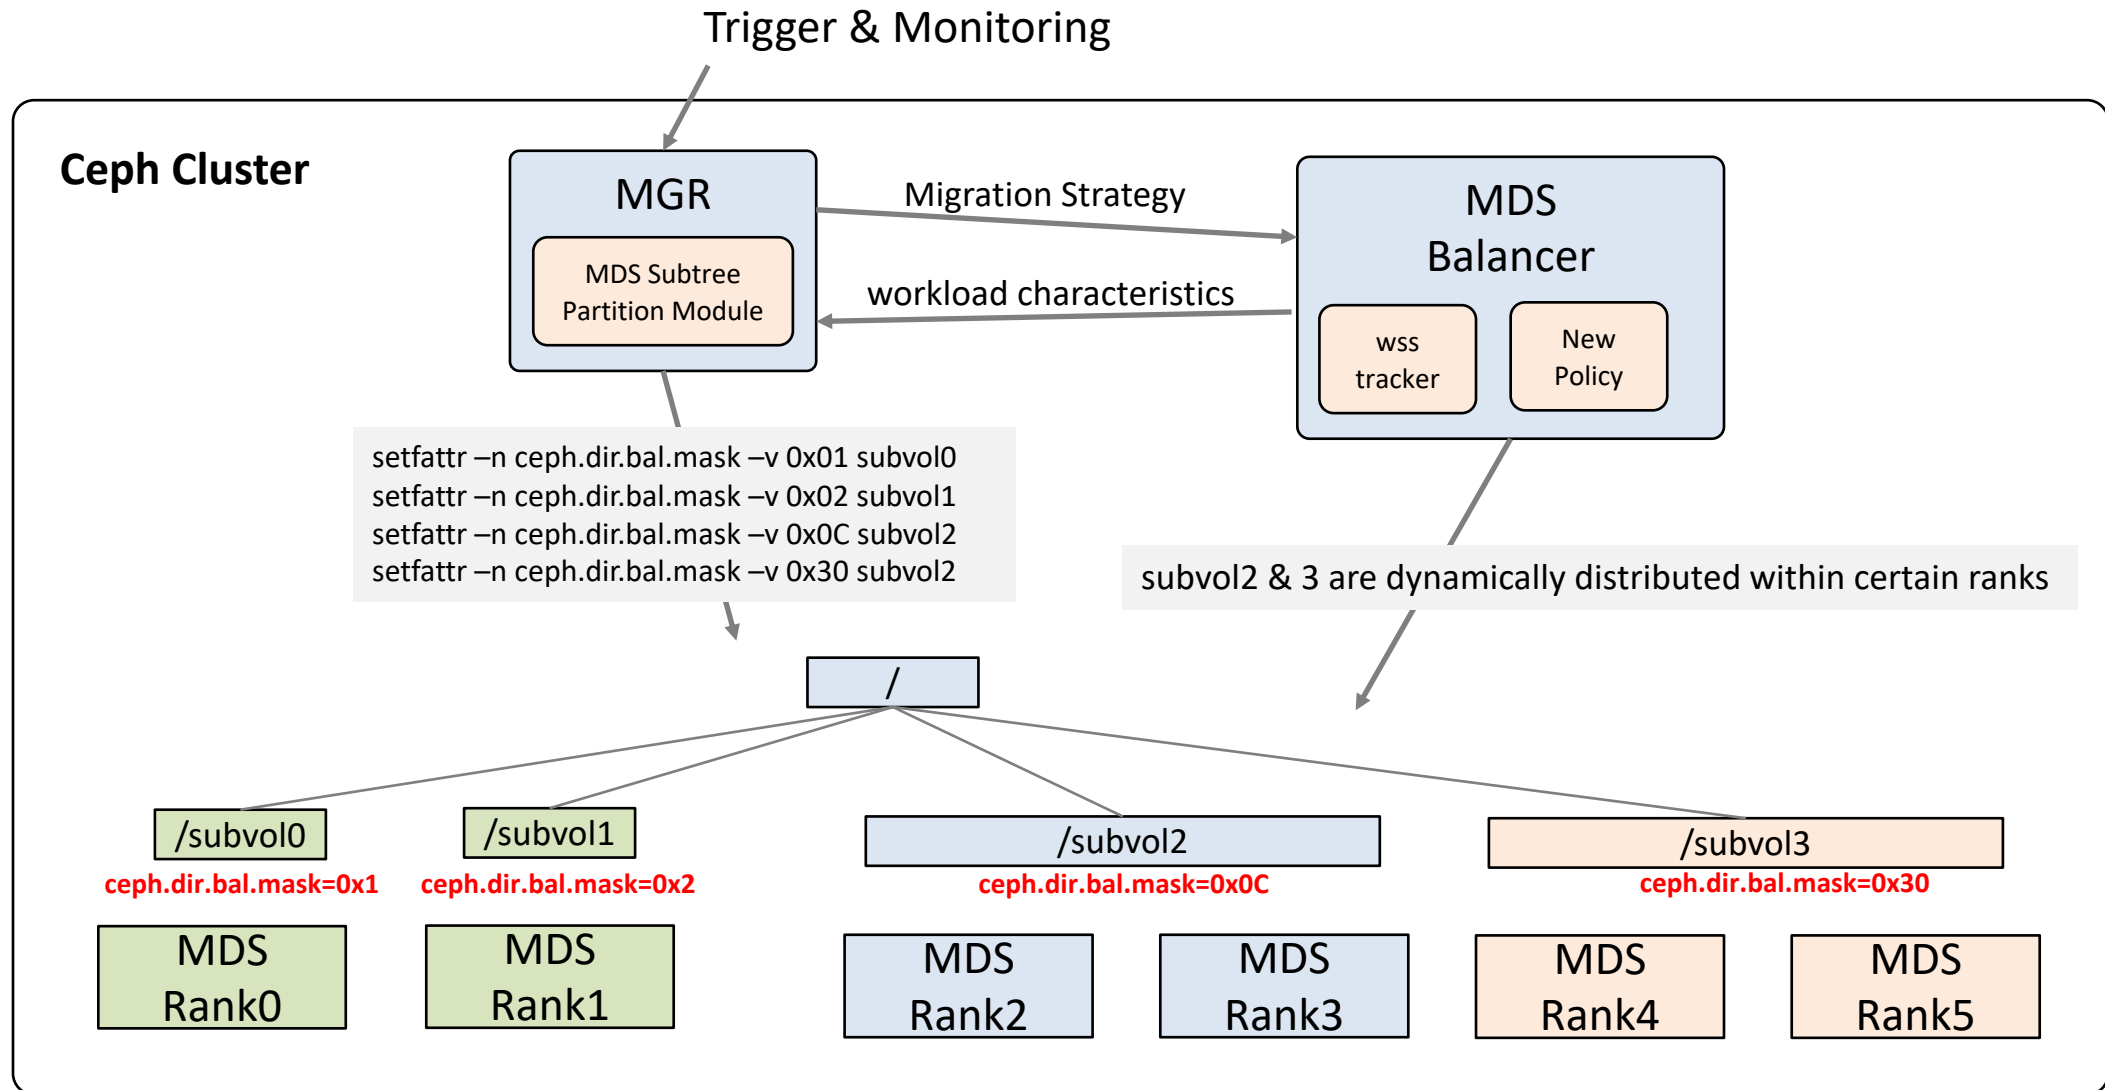
inodes are balanced and dir fetch count decreases and performance increases



# Outline of Contents

- Introduction and background
- CephFS subtree partitioning
- Combining static and dynamic partitioning schemes
  - Workload based static partitioner with `bal_rank_mask`
  - A new CephFS MDS balancer with `ceph.dir.bal.mask`
- Future works and conclusions

# A New MDS Balancer with ceph.dir.bal.mask



# Implementation

- rank mask option per subdir as a virtual extended attribute
  - A target subdir is dynamically within certain MDS ranks (e.g., rank0 and 1)

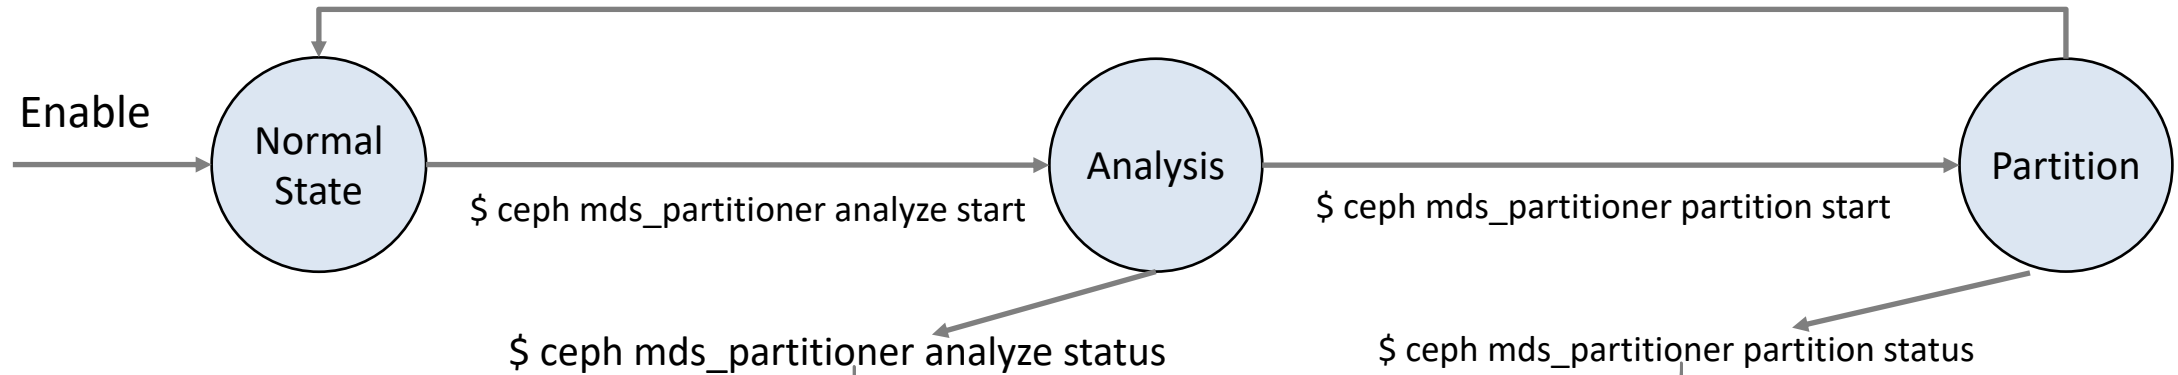
```
setfattr -n ceph.dir.bal.mask -v 0x3 /cephfs/home/yongseok
```

- MDS Subtree Partition Module in MGR

```
ceph mgr module enable mds_partitioner
ceph mds_partitioner analyze start # analyze client workloads obtained from MDSs
ceph mds_partitioner analyze status # report analysis results and recommend optimal the number of MDSs
ceph mds_partitioner partition start # start partitioning
ceph mds_partitioner partition status # report partitioning status
```

- MDS Balancer modifications
  - Working set size tracker
  - Migrate subdirs based on ceph.dir.bal.mask values of subdirs
  - Minimize MDS slow requests

# Example of Operation Flow



Name	wss	reqs	rentries	workload	Current Ranks	New Ranks	Migration Progress	Migration Status
Subvol1	1,000,000	1,203,030	50,000,000	339	0	0,1	20%	In-progress
Subvol2	700,000	500,000	1,000,000	137	1	2	0%	Ready
subvol3	100,000	5,000	5,000,000	21	2	2	100%	Done
subvol4	3,000	20,000	70,000	3	2	2	100%	Done
Total	1,803,000	1,728,030	56,070,000	500				

wss: working set size  
 reqs: requests  
 rentries: files + dirs

**workload** = (working\_set\_size / total\_working\_set \* 2  
 + requests / total\_requests \* 2  
 + rentries / total\_rentries) \* 100



# Conclusions

- We employ CephFS as a shared file service in LINE's cloud
- We compared static and dynamic partitioning schemes
  - Dynamic partitioning incurs metadata migration cost
  - Static partitioning present uneven workload distribution
- We presented how to combine both static and dynamic partitioning scheme effectively
- We will contribute our work on a new partitioning to the community

# **Thank you!**

Any Questions?

# References

- Ceph: A Scalable, High-Performance Distributed File System, Sage Weil, OSDI'06
- Overview and Status of the Ceph File System, Patrick Donnelly, 2018, <https://indico.cern.ch/event/644915/>
- ceph-linode for CephFS testing, Patrick Donnelly, <https://github.com/batrick/ceph-linode>
- CephFS with OpenStack Manila based on BlueStore and Erasure Code, 2018, <https://cutt.ly/dc7Qnn7>
- Revisiting CephFS MDS and mClock QoS Scheduler, Yongseok Oh, 2021, <https://cutt.ly/A4s8AsC>
- Ephemeral Pinning: A Dynamic Metadata Management Strategy for CephFS, Sidharth Anupkrishnan, 2020 [https://www.youtube.com/watch?v=zimAEm\\_8efA](https://www.youtube.com/watch?v=zimAEm_8efA)
- Optimizing CephFS with Combining MDS QoS Scheduling and Static-Dynamic Subtree Partitioning, Yongseok Oh, <https://ceph2023.sched.com/event/1JKas/optimizing-cephfs-with-combining-mds-qos-scheduling-and-static-dynamic-subtree-partitioning-yongseok-oh-jinmyeong-lee-line>

# VDBench Parameter Configs

client 01~02

```
fsd=fsd1,anchor=/mnt,depth=1,width=1,files=50000,size=4k  
fwd=fwd1,fsd=fsd1,operation=getattr,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
fwd=fwd2,fsd=fsd1,operation=read,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
rd=rd1,fwd=(fwd1,fwd2),fwdrate=max,elapsed=86400,format=yes,interval=10
```

client 03~20

```
fsd=fsd1,anchor=/mnt,depth=1,width=1,files=50000,size=4k  
fwd=fwd1,fsd=fsd1,operation=getattr,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
fwd=fwd2,fsd=fsd1,operation=create,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
fwd=fwd3,fsd=fsd1,operation=delete,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
fwd=fwd4,fsd=fsd1,operation=write,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
rd=rd1,fwd=(fwd1,fwd2,fwd3,fwd4),fwdrate=max,elapsed=86400,format=yes,interval=1
```

client 21~22

```
fsd=fsd1,anchor=/mnt,depth=1,width=10,files=50000,size=4k  
fwd=fwd1,fsd=fsd1,operation=getattr,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
fwd=fwd2,fsd=fsd1,operation=read,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
rd=rd1,fwd=(fwd1,fwd2),fwdrate=max,elapsed=86400,format=yes,interval=10
```

client 23~40

```
fsd=fsd1,anchor=/mnt,depth=1,width=10,files=50000,size=4k  
fwd=fwd1,fsd=fsd1,operation=getattr,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
fwd=fwd2,fsd=fsd1,operation=create,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
fwd=fwd3,fsd=fsd1,operation=delete,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
fwd=fwd4,fsd=fsd1,operation=write,xfersize=4k,fileio=sequential,fileselect=random,threads=1  
rd=rd1,fwd=(fwd1,fwd2,fwd3,fwd4),fwdrate=max,elapsed=86400,format=yes,interval=1
```


client 41

```
fsd=fsd1,anchor=/mnt,depth=1,width=100,files=50000,size=4k  
fwd=fwd1,fsd=fsd1,operation=getattr,xfersize=4k,fileio=sequential,fileselect=random,threads=16  
fwd=fwd2,fsd=fsd1,operation=read,xfersize=4k,fileio=sequential,fileselect=random,threads=16  
rd=rd1,fwd=(fwd1,fwd2),fwdrate=max,elapsed=86400,format=yes,interval=10
```

client 42

```
fsd=fsd1,anchor=/mnt,depth=1,width=100,files=50000,size=4k  
fwd=fwd1,fsd=fsd1,operation=getattr,xfersize=4k,fileio=sequential,fileselect=random,threads=16  
fwd=fwd2,fsd=fsd1,operation=create,xfersize=4k,fileio=sequential,fileselect=random,threads=16  
fwd=fwd3,fsd=fsd1,operation=delete,xfersize=4k,fileio=sequential,fileselect=random,threads=16  
fwd=fwd4,fsd=fsd1,operation=write,xfersize=4k,fileio=sequential,fileselect=random,threads=16  
rd=rd1,fwd=(fwd1,fwd2,fwd3,fwd4),fwdrate=max,elapsed=86400,format=yes,interval=10
```

# Ceph CI/CD System in LINE




## Ceph Test Frameworks

[teuthology]

[verdology]

The screenshot shows the Pulpito interface for Teuthology. It displays a table of test results for a specific job. The table includes columns for Status, Job ID, Links, Posted, Started, Updated, Runtime, Duration, In Waiting, Machine, Teuthology Branch, OS Type, OS Version, and Nodes. The test results show a mix of passed and failed jobs.

Status	Job ID	Links	Posted	Started	Updated	Runtime	Duration	In Waiting	Machine	Teuthology Branch	OS Type	OS Version	Nodes
pass	187		2021-04-09 01:41:22	2021-04-09 03:12:41	2021-04-09 03:40:40	0:27:59	0:25:31	0:02:28	vps	master	centos	7.8	2
pass	188		2021-04-09 01:41:23	2021-04-09 03:12:43	2021-04-09 03:34:42	0:21:59	0:18:42	0:03:17	vps	master	centos	7.8	2
pass	189		2021-04-09 01:41:24	2021-04-09 03:34:44	2021-04-09 04:08:43	0:33:59	0:31:56	0:02:03	vps	master	centos	7.8	2
pass	190		2021-04-09 01:41:25	2021-04-09 03:40:42	2021-04-09 04:02:41	0:21:59	0:18:34	0:03:25	vps	master	centos	7.8	2
pass	191		2021-04-09 01:41:26	2021-04-09 04:02:43	2021-04-09 04:22:42	0:19:59	0:17:21	0:02:38	vps	master	centos	7.8	2
pass	192		2021-04-09 01:41:27	2021-04-09 04:08:45	2021-04-09 04:28:44	0:19:59	0:17:27	0:02:32	vps	master	centos	7.8	2
pass	193		2021-04-09 01:41:28	2021-04-09 04:22:44	2021-04-09 04:54:43	0:31:59	0:28:59	0:03:00	vps	master	centos	7.8	2
pass	194		2021-04-09 01:41:29	2021-04-09 04:28:46	2021-04-09 05:00:45	0:31:59	0:29:42	0:02:17	vps	master	centos	7.8	2
pass	195		2021-04-09 01:41:30	2021-04-09 04:54:45	2021-04-09 05:32:44	0:37:59	0:34:57	0:03:02	vps	master	centos	7.8	2
pass	196		2021-04-09 01:41:31	2021-04-09 05:00:47	2021-04-09 05:22:47	0:22:00	0:19:46	0:02:14	vps	master	centos	7.8	2
pass	197		2021-04-09 01:41:32	2021-04-09 05:22:49	2021-04-09 05:50:48	0:27:59	0:25:25	0:02:34	vps	master	centos	7.8	2
pass	198		2021-04-09 01:41:32	2021-04-09 05:32:46	2021-04-09 05:54:46	0:22:00	0:19:06	0:02:54	vps	master	centos	7.8	2
pass	199		2021-04-09 01:41:33	2021-04-09 05:50:50	2021-04-09 06:12:50	0:22:00	0:18:28	0:03:32	vps	master	centos	7.8	2
pass	200		2021-04-09 01:41:34	2021-04-09 05:54:48	2021-04-09 06:24:47	0:29:59	0:27:50	0:02:09	vps	master	centos	7.8	2



# CephFS Test Flow of Verdology

